

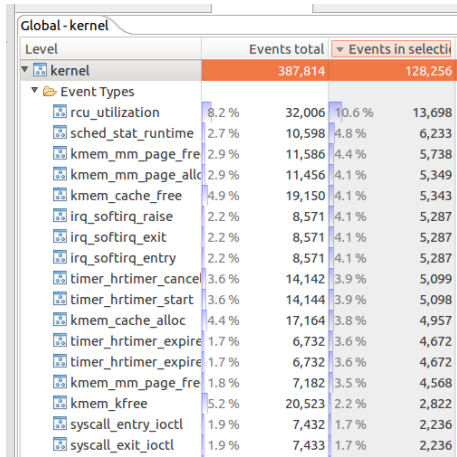
# Assignment Report

Isaac Souza

April 25, 2017

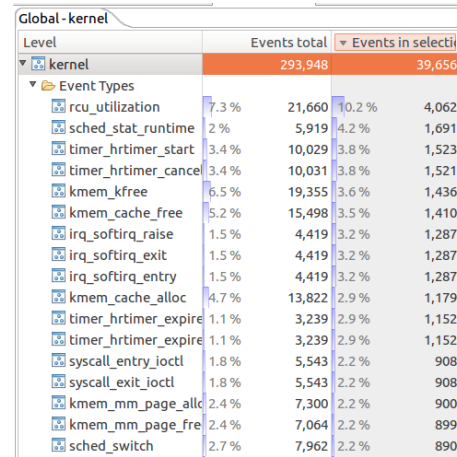
This document is a report of a tracing analysis that I have made in order to demonstrate my familiarity with tracing. I used a code that I have written to calculate the approximate value of  $\pi$  using the Archimedes' Method. First, I used a serial version of the program written in C language. Setting the number of iterations to 100,000,000, I executed the program while using *lttng* to trace the system kernel. Then, I used *trace-compass* to measure the elapsed time (around 2 seconds) and how it interacted with the system kernel. Then, using OpenMP, I made the main loop of the program parallel. Again, using *lttng* and *trace-compass*, I executed the parallelized version of the program and found out that the performance was worst then the serial version (around 4 seconds), different from the expected behavior. In order to find the problem, I imported the two traces on *trace-compass*. Inspecting the parallelized version trace, I saw that each thread on the parallelized version was taking around 4 seconds to complete its job, showing that the problem was with my parallelization.

After some quick research and trace analysis, I found out that the problem with my code resided in the fact that each thread was calling the `drand48()` function which is not thread-safe, since it can make the process become locked until another thread finish reading the RNG state data memory address. The figure 1a shows the event count on the time period of the program execution. To address the problem, I used the `srand48_r()` and `drand48_r()` functions instead of the former ones. As a result, this second version of the parallelized program ran faster (around 1 second) then the serial one, as expected. The figure 1b shows the event count for this second version of the code. Comparing the two results on *trace-compass*, I saw that, on the first version, the percentage of events of the type `kmem_page_alloc_zone_locked` (fourth on its list) is bigger than the second one, demonstrating that each thread stays more time locked.



Level	Events total	Events in selecti
kernel	387,814	128,256
Event Types		
rcu_utilization	8.2 %	32,006
sched_stat_runtime	2.7 %	10,598
kmem_mm_page_fre	2.9 %	11,586
kmem_mm_page_allc	2.9 %	11,456
kmem_cache_free	4.9 %	19,150
irq_softirq_raise	2.2 %	8,571
irq_softirq_exit	2.2 %	8,571
irq_softirq_entry	2.2 %	8,571
timer_hrtimer_cancel	3.6 %	14,142
timer_hrtimer_start	3.6 %	14,144
kmem_cache_alloc	4.4 %	17,164
timer_hrtimer_expire	1.7 %	6,732
timer_hrtimer_expire	1.7 %	6,732
kmem_mm_page_fre	1.8 %	7,182
kmem_kfree	5.2 %	20,523
syscall_entry_ioctl	1.9 %	7,432
syscall_exit_ioctl	1.9 %	7,433

(a) First version of the program.



Level	Events total	Events in selecti
kernel	293,948	39,656
Event Types		
rcu_utilization	7.3 %	21,660
sched_stat_runtime	2 %	5,919
timer_hrtimer_start	3.4 %	10,029
timer_hrtimer_cancel	3.4 %	10,031
kmem_kfree	6.5 %	19,355
kmem_cache_free	5.2 %	15,498
irq_softirq_raise	1.5 %	4,419
irq_softirq_exit	1.5 %	4,419
irq_softirq_entry	1.5 %	4,419
kmem_cache_alloc	4.7 %	13,822
timer_hrtimer_expire	1.1 %	3,239
timer_hrtimer_expire	1.1 %	3,239
syscall_entry_ioctl	1.8 %	5,543
syscall_exit_ioctl	1.8 %	5,543
kmem_mm_page_alloc	2.4 %	7,300
kmem_mm_page_fre	2.4 %	7,064
sched_switch	2.7 %	7,962

(b) Second version of the program.

Figure 1: The event count for the two version of the parallelized program.